

Voice-Enabled Government Scheme Recommendation Agent

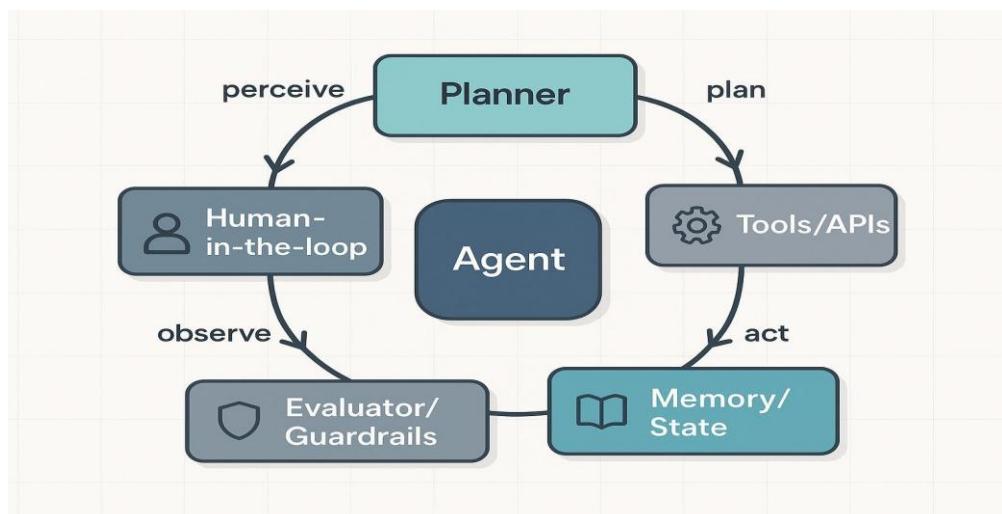
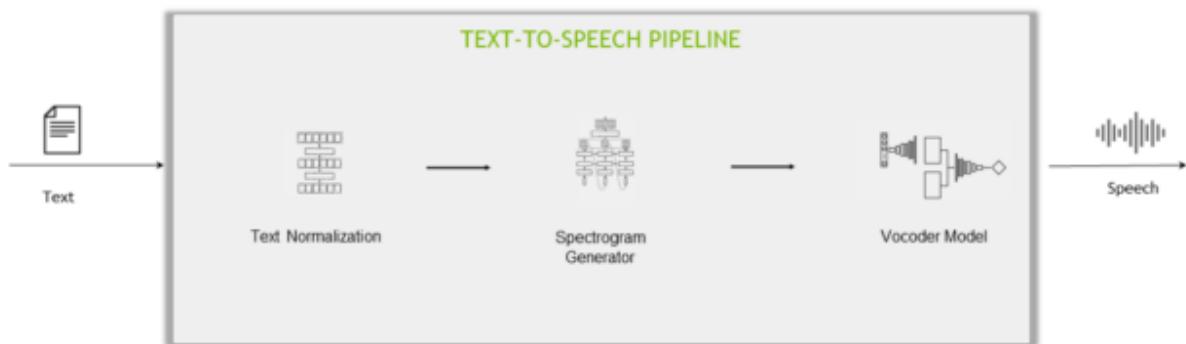
1. System Overview

The system is a **voice-driven AI agent** that interacts with a user, collects required personal information step by step, confirms inputs to reduce speech errors, and determines eligibility for government welfare schemes.

Core goals:

- Voice-first interaction (STT + TTS)
- Robust handling of speech errors
- Deterministic decision-making
- Transparent eligibility reasoning

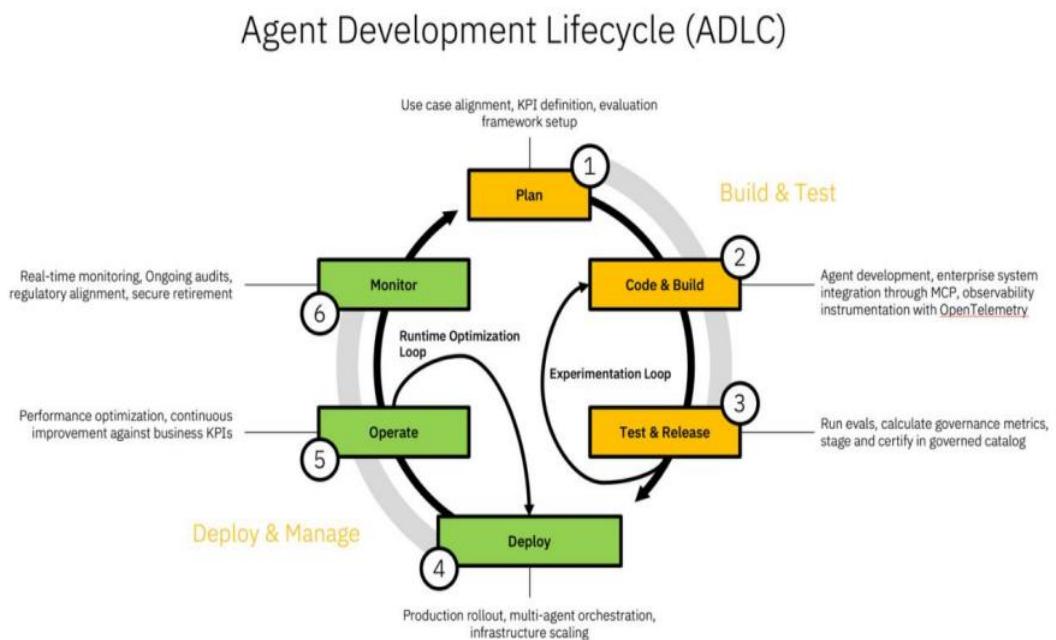
2. High-Level Architecture

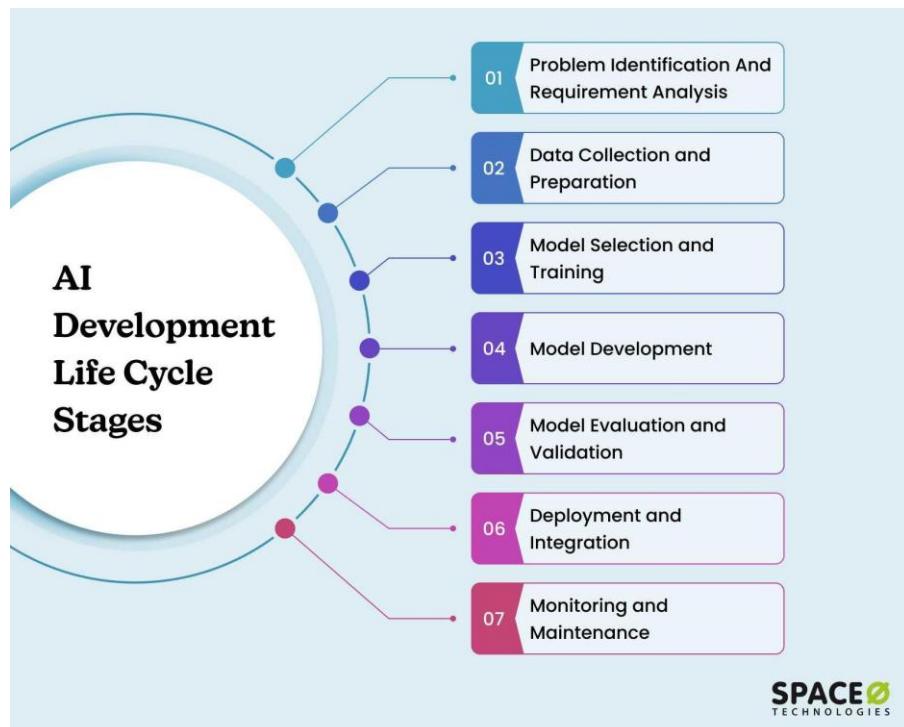


Components

Layer	Component	Responsibility
Interface	STT (Whisper)	Converts user speech → text
Interface	TTS (gTTS)	Converts agent text → speech
Agent Core	Planner	Decides next action
Agent Core	Executor	Executes the decided action
Agent Core	Evaluator	Validates output & decides state
Memory	Agent Memory	Stores user profile + state
Tools	Scheme Retriever	Loads scheme data
Tools	Eligibility Checker	Matches user vs scheme rules

3. Agent Lifecycle



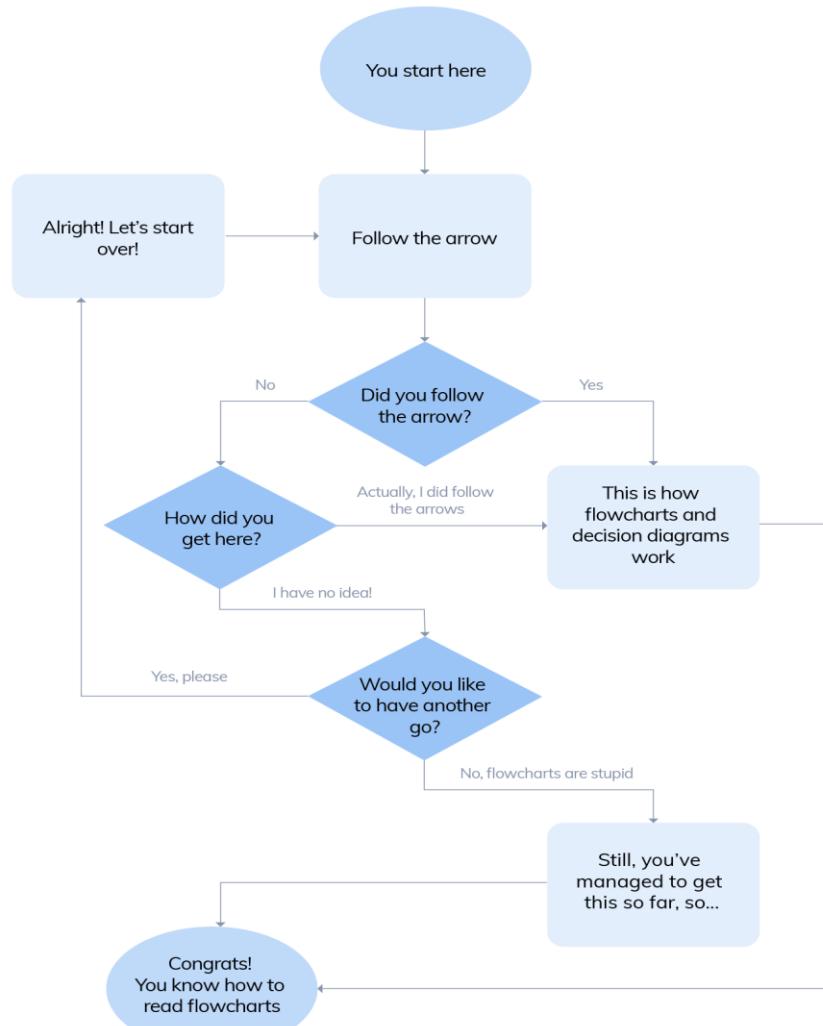
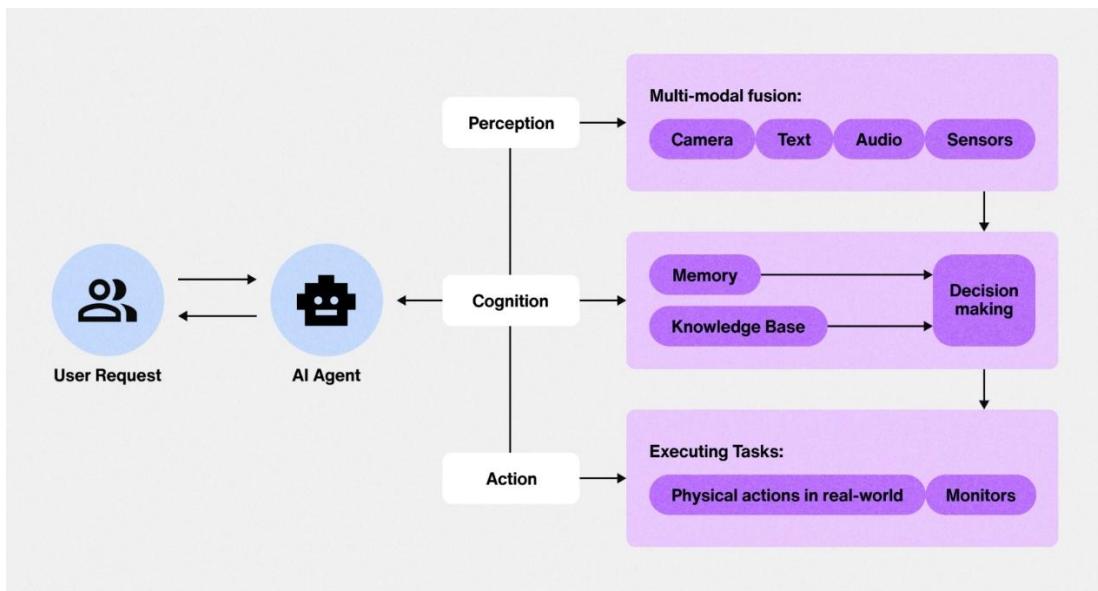


Lifecycle Steps

1. Greeting
2. User speaks
3. Speech → Text (STT)
4. Confirmation loop
5. Planner decision
6. Executor action
7. Evaluator validation
8. Repeat OR Complete

This loop continues until all required information is collected and eligibility is computed.

4. Detailed Decision Flow



Decision Logic (Planner)

The planner always selects **exactly one** action:

IF contradictions exist

→ HANDLE_CONTRADICTION

ELSE IF any required field missing

→ ASK_MISSING_INFO:<field>

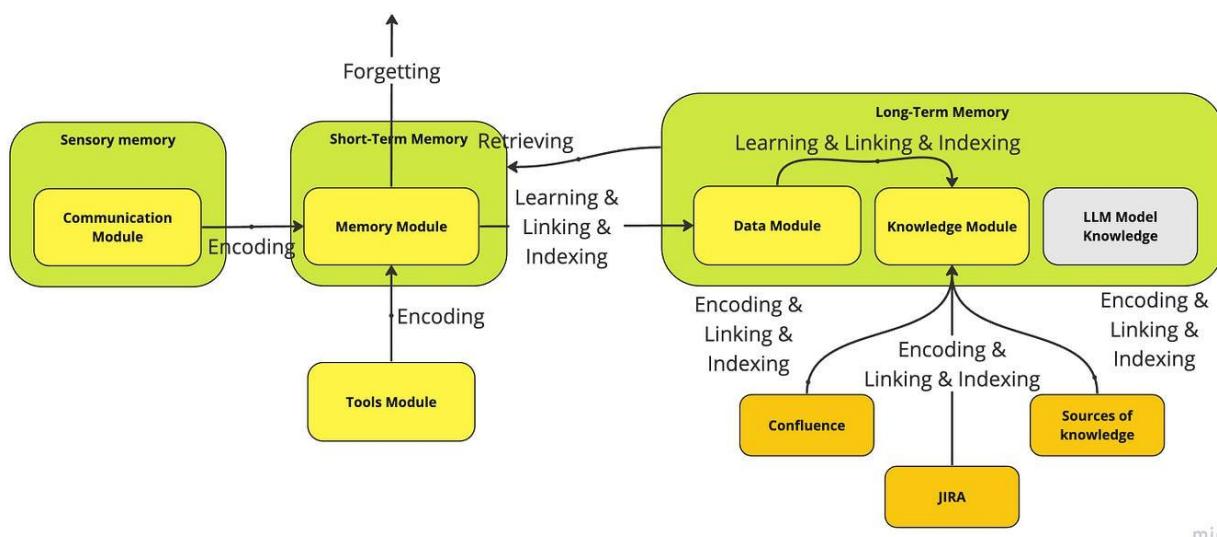
ELSE

→ CHECK_ELIGIBILITY

This guarantees:

- Deterministic behavior
- No hallucinated actions
- Predictable execution

5. Agent Memory Design



Memory Structure

```
{  
  "name": "string",  
  "age": "int",
```

```
"income": "int",
"gender": "male/female",
"bpl": "boolean",
"housing_status": "homeless/owned/rented"
}
```

Additional Memory State

- pending_confirmation → for ASR correction
- contradictions[] → conflict tracking
- history[] → conversation log
- state → collecting_info / completed

6. Prompt Architecture (Planner)

Planner Prompt (Design Principle)

The planner never generates user-facing text.
It only returns **one of three commands**.

Prompt Template

You are an internal planner for an AI agent.

User text:

{user_text}

Current memory:

{memory}

Contradictions:

{contradictions}

Choose exactly ONE:

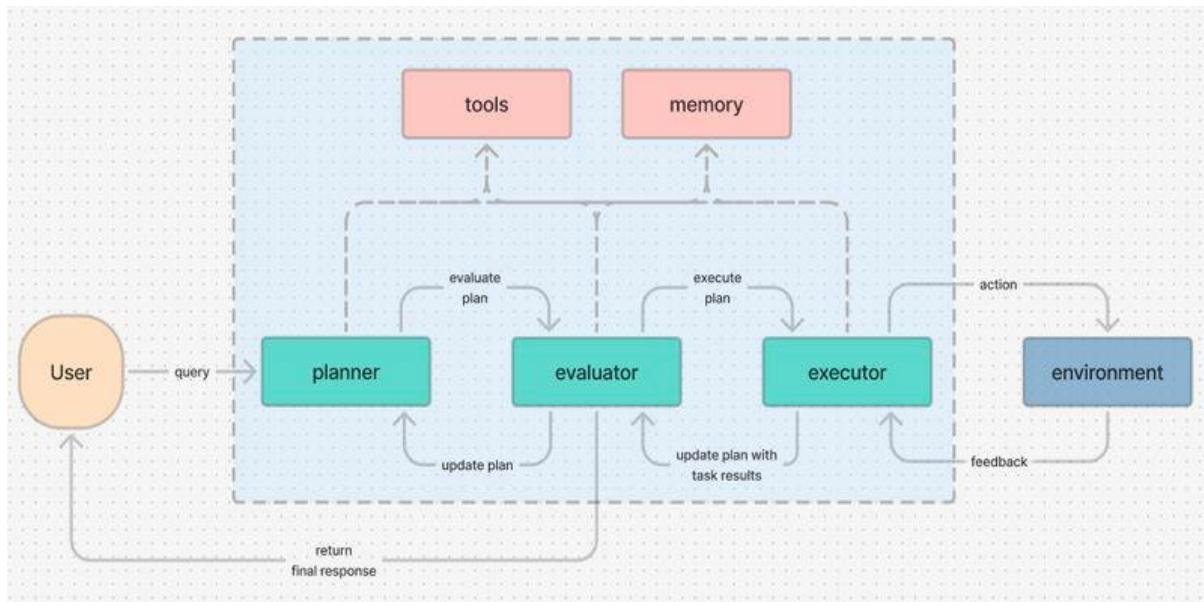
1. ASK_MISSING_INFO:<field>
2. CHECK_ELIGIBILITY

3. HANDLE CONTRADICTION

Why This Works

- Prevents hallucinations
- Keeps logic auditable
- Makes agent explainable

7. Executor–Evaluator Pattern



Executor

- Translates planner action → real operation
- Calls tools
- Fetches schemes
- Checks eligibility

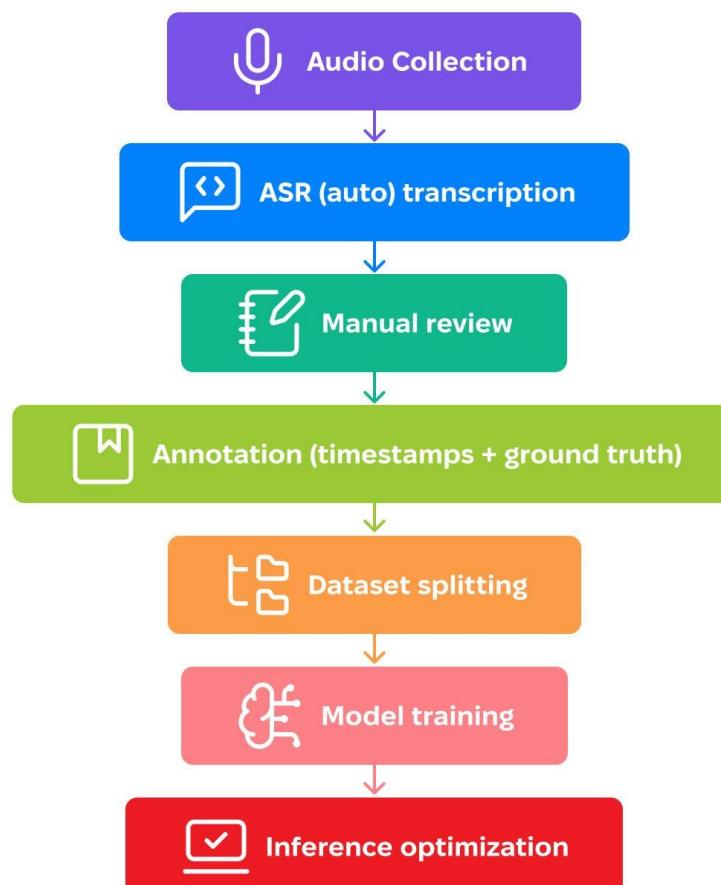
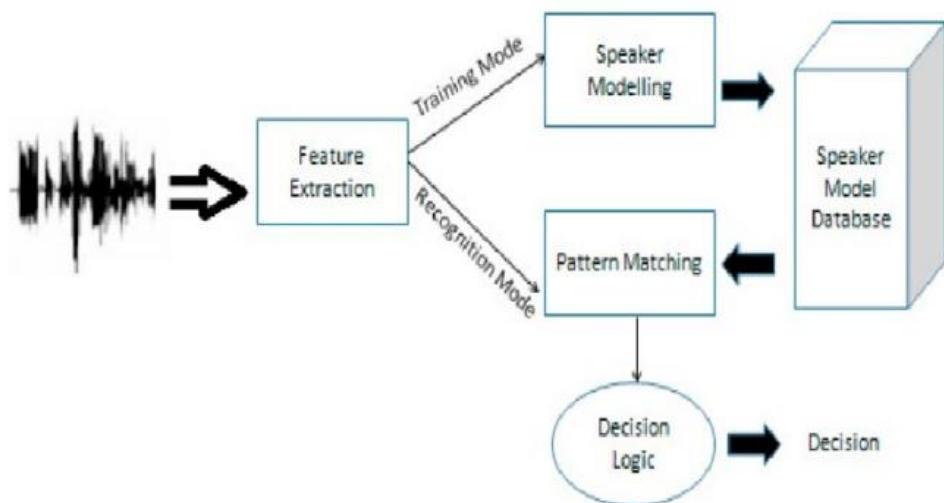
Evaluator

- Validates executor output
- Decides:
 - WAITING_FOR_USER
 - COMPLETED
 - FAILED

This separation mirrors **production-grade agent systems**.

8. Speech Error Handling (Key Design Strength)

Confirmation-Based Voice Pattern



Flow:

1. User speaks
2. STT transcribes
3. Agent repeats text
4. User confirms (Yes / No)
5. Only confirmed data is stored

Benefit

- Works even with noisy STT
- Transparent to user
- Highly reliable for real-world usage

9. Scheme Eligibility Engine

Rule-Based Matching

Each scheme defines explicit criteria:

```
{  
  "income_max": 300000,  
  "housing_status": "homeless"  
}
```

Eligibility is checked using **deterministic rules**, not LLM guesses.

Advantages

- Explainable decisions
- No hallucination risk
- Easy to extend

10. End-to-End Flow Summary

User Voice

↓

Speech-to-Text (Whisper)

↓

Confirmation Loop

↓
Planner (LLM – decision only)
↓
Executor (tools + rules)
↓
Evaluator (state decision)
↓
Text-to-Speech (Odia response)